

## Kubernetes 05

1. 1 . Create a namespace dev-environment and apply a resource-based quota that restricts the number of pods to 3 and services to 2.

- Create dev environment
- kubectl create namespace dev-environment

```
[ec2-user@ip-172-31-4-234 ~]$ kubectl create namespace dev-environment
namespace/dev-environment created
[ec2-user@ip-172-31-4-234 ~]$ kubectl get namespaces
NAME          STATUS   AGE
default       Active   9d
dev-environment Active   8s
kube-flannel  Active   9d
kube-node-lease Active   9d
kube-public   Active   9d
kube-system   Active   9d
```

- Now create a yaml file
- Vi resource-quota.yaml
- apiVersion: v1
- kind: ResourceQuota
- metadata:
- name: dev-quota
- namespace: dev-environment
- spec:
- hard:
- pods: "3"
- services: "2"

- Add deploy it

```
default          Active  9d
dev-environment Active  8s
kube-flannel    Active  9d
kube-node-lease Active  9d
kube-public     Active  9d
kube-system     Active  9d
[ec2-user@ip-172-31-4-234 ~]$ vi resource-quota.yaml
[ec2-user@ip-172-31-4-234 ~]$ kubectl apply -f resource-quota.yaml
resourcequota/dev-quota created
```

- Verify
- `kubectl get resourcequota -n dev-environment`

```
 kube-system          Active  9d
[ec2-user@ip-172-31-4-234 ~]$ vi resource-quota.yaml
[ec2-user@ip-172-31-4-234 ~]$ kubectl apply -f resource-quota.yaml
resourcequota/dev-quota created
[ec2-user@ip-172-31-4-234 ~]$ kubectl get resourcequota -n dev-environment
NAME      AGE      REQUEST           LIMIT
dev-quota  8s      pods: 0/3, services: 0/2
[ec2-user@ip-172-31-4-234 ~]$ █
```

- 2 . Create a pod in the prod-environment namespace with 0.2 CPU and 200Mi memory requests, and 0.5 CPU and 500Mi memory limits.

- Create prod environment
- `kubectl create namespace prod-environment`
- now create a yaml file

```
prod-pod  1/1    Running   0          11s
[ec2-user@ip-172-31-4-234 ~]$ cat prod-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: prod-pod
  namespace: prod-environment
spec:
  containers:
    - name: nginx-container
      image: nginx
      resources:
        requests:
          cpu: "200m"
          memory: "200Mi"
        limits:
          cpu: "500m"
          memory: "500Mi"
```

```
[ec2-user@ip-172-31-4-234 ~]$ █
```

- 
- Deploy it
- Kubectl apply -f prod-pod.yaml
- Verify it
- kubectl get pods -n prod-environment

```
[ec2-user@ip-172-31-4-234 ~]$ kubectl apply -f prod-pod.yaml
pod/prod-pod created
[ec2-user@ip-172-31-4-234 ~]$ kubectl get pods -n prod-environment
NAME      READY   STATUS    RESTARTS   AGE
prod-pod  1/1     Running   0          11s
```

- kubectl describe pod prod-pod -n prod-environment
- this show the memory and cpu
  - Guaranteed **0.2 CPU / 200Mi memory**
  - Capped at **0.5 CPU / 500Mi memory**
-

```

        Port: <none>
        Host Port: <none>
        State: Running
          Started: Mon, 05 Jan 2026 05:07:28 +0000
          Ready: True
        Restart Count: 0
        Limits:
          cpu: 500m
          memory: 500Mi
        Requests:
          cpu: 200m
          memory: 200Mi
        Environment: <none>
        Mounts:
          /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-g2nrj (ro)
Conditions:
  Type          Status
  PodReadyToStartContainers True
  Initialized   True
  Ready         True
  ContainersReady True
  PodScheduled  True
Volumes:
  kube-api-access-g2nrj:
    Type:      Projected (a volume that contains injected data from multiple source
●

```

3 . In the staging-environment namespace, set a LimitRange that assigns default CPU and memory limits (300m CPU, 600Mi memory) and applies a minimum and maximum CPU.

- Create Environmetn
- Kubectl creaste namespace staging

```

        Normal  Started  2m56s  kubelet            started container nginx-container
[ec2-user@ip-172-31-4-234 ~]$ kubectl create namespace staging-environment
namespace/staging-environment created
[ec2-user@ip-172-31-4-234 ~]$ kubectl get namespaces
NAME           STATUS  AGE
default        Active  9d
dev-environment Active  11m
kube-flannel   Active  9d
kube-node-lease Active  9d
kube-public    Active  9d
kube-system    Active  9d
prod-environment Active  5m34s
●  staging-environment  Active  6s

```

- Now,
- Write a limitrange.yml file
- apiVersion: v1
- kind: LimitRange
- metadata:

- name: staging-limits
- namespace: staging-environment
- spec:
- limits:
  - type: Container
- default:
  - cpu: "300m"
  - memory: "600Mi"
- defaultRequest:
  - cpu: "300m"
  - memory: "600Mi"
- min:
  - cpu: "100m"
- max:
  - cpu: "500m"
- Now deploy it
- Kubectl apply -f limitrange.yml

```
[ec2-user@ip-172-31-4-234 ~]$ kubectl apply -f limitrange.yml
limitrange/staging-limits created
[ec2-user@ip-172-31-4-234 ~]$ kubectl get limitrange -n staging-environment
NAME          CREATED AT
staging-limits 2026-01-05T05:12:47Z
```

- Now verify it
- kubectl get limitrange -n staging-environment
- get detailed view
- kubectl describe limitrange staging-limits -n staging-environment

```

[ec2-user@ip-172-31-4-234 ~]$ kubectl describe limitrange staging-limits -n staging-environment
Name:      staging-limits
Namespace: staging-environment
Type       Resource   Min   Max   Default Request  Default Limit  Max Limit/Request Ratio
----       -----   ---   ---   -----          -----        -----
Container  memory     -     -    600Mi           600Mi         -
Container  cpu        100m  500m  300m           300m         -
[ec2-user@ip-172-31-4-234 ~]$ █

```

- 
- 4 . Create a pod and a NodePort service in the default namespace, then create another pod in the test namespace and communicate between them using Service DNS.

- Vi default-pod.yaml
  - apiVersion: v1
  - kind: Pod
  - metadata:
  - name: web-pod
  - labels:
  - app: web
  - spec:
  - containers:
  - - name: nginx
  - image: nginx
  - ports:
  - - containerPort: 80
  - Deploy it
  - Kubectl apply -f default-pod.yml
- ```

[ec2-user@ip-172-31-4-234 ~]$ vi default-pod.yaml
[ec2-user@ip-172-31-4-234 ~]$ kubectl apply -f de
default-pod.yaml deployment.yaml
[ec2-user@ip-172-31-4-234 ~]$ kubectl apply -f default-pod.yaml
● pod/web-pod created

```
- Create a NodePort Service in default

- Vi web.service.yml
- apiVersion: v1
- kind: Service
- metadata:
- name: web-service
- spec:
- type: NodePort
- selector:
- app: web
- ports:
- - port: 80
- targetPort: 80
- nodePort: 30080
- Deploy it
- kubectl apply -f web-service.yaml

```
[ec2-user@ip-172-31-4-234 ~]$ kubectl get endpoints recreate-service
NAME           ENDPOINTS   AGE
recreate-service  10.244.1.35:80,10.244.2.42:80,10.244.2.44:80  5d22h
[ec2-user@ip-172-31-4-234 ~]$ kubectl describe svc recreate-service
Name:            recreate-service
Namespace:       default
Labels:          <none>
Annotations:    <none>
Selector:        app=recreate-app
Type:            NodePort
IP Family Policy: SingleStack
IP Families:    IPv4
IP:              10.102.74.118
IPs:             10.102.74.118
Port:            <unset>  80/TCP
TargetPort:      80/TCP
NodePort:        <unset>  32463/TCP
Endpoints:      10.244.1.35:80,10.244.2.42:80,10.244.2.44:80
Session Affinity: None
External Traffic Policy: Cluster
Events:          <none>
```

- Dns is working perfect

```
kube-scheduler-ip-172-31-4-234.us-west-1.compute.internal      1/1   Ru
[ec2-user@ip-172-31-4-234 ~]$ kubectl exec -it client-pod -n test -- sh
/ # nslookup recreate-service.default.svc.cluster.local
Server:          10.96.0.10
Address:         10.96.0.10:53

Name:    recreate-service.default.svc.cluster.local
Address: 10.102.74.118

/ #
/ #
```

- 

1. 5 . Apply a LimitRange with a max limit/request ratio of 2 for memory in the performance-environment namespace, and test by creating a pod with mismatched resource requests and limits.

- Create a namespace
- Now write a Yaml file
- apiVersion: v1
- kind: LimitRange
- metadata:
- name: memory-ratio-limit
- namespace: performance-environment
- spec:
- limits:
- - type: Container
- maxLimitRequestRatio:
- memory: "2"
- Deploy it
- By
- Kubectl apply -f memory-limitrange.yml

```
[ec2-user@ip-172-31-4-234 ~]$ kubectl create namespace performance-environment  
namespace/performance-environment created  
[ec2-user@ip-172-31-4-234 ~]$ vi memory-limitrange.yaml  
[ec2-user@ip-172-31-4-234 ~]$ kubectl apply -f memory-limitrange.yaml  
limitrange/memory-ratio-limit created
```

- Now
- Write a testing yaml file as
- Bad-pod.yaml
- apiVersion: v1
- kind: Pod
- metadata:
- name: bad-memory-pod
- namespace: performance-environment
- spec:
- containers:
- - name: test
- image: nginx
- resources:
- requests:
- memory: "200Mi"
- limits:
- memory: "500Mi" when we try to deploy it

```
[ec2-user@ip-172-31-4-234 ~]$ kubectl apply -f bad-pod.yaml
Error from server (Forbidden): error when creating "bad-pod.yaml": pods
ed ratio is 2.500000
[ec2-user@ip-172-31-4-234 ~]$ cat bad-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: bad-memory-pod
  namespace: performance-environment
spec:
  containers:
    - name: test
      image: nginx
      resources:
        requests:
          memory: "200Mi"
        limits:
          memory: "500Mi"
```

- [ec2-user@ip-172-31-4-234 ~]\$ █

- It doesn't deploy
- Kubectl apply -f bad-pod.yaml
- It throws error

```
[ec2-user@ip-172-31-4-234 ~]$ vi bad-pod.yaml
[ec2-user@ip-172-31-4-234 ~]$ kubectl apply -f bad-pod.yaml
Error from server (Forbidden): error when creating "bad-pod.yaml": pods "bad-memory-pod" is forbidden: memory max limit to request ratio per container is 2, but provided ratio is 2.500000
[ec2-user@ip-172-31-4-234 ~]$ cat bad-pod.yaml
```